

Tidy data & dplyr

Lecture 06

Dr. Colin Rundel



Tidy data

country	year	cases	population
Afghanistan	1999	37745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	213766	1280425583

variables

country	year	cases	population
Afghanistan	1999	37745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	213766	1280425583

observations

country	year	cases	population
Afghanistan	1999	37745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	213766	1280425583

values

Tidy vs Untidy

Happy families are all alike; every unhappy family is unhappy in its own way

— Leo Tolstoy, Anna Karenina

```
# A tibble: 317 × 7
```

	artist	track	date.entered	wk1	wk2	wk3	wk4
	<chr>	<chr>	<date>	<dbl>	<dbl>	<dbl>	<dbl>
1	2 Pac	Baby Don't Cry...	2000-02-26	87	82	72	77
2	2Ge+her	The Hardest Pa...	2000-09-02	91	87	92	NA
3	3 Doors Down	Kryptonite	2000-04-08	81	70	68	67
4	3 Doors Down	Loser	2000-10-21	76	76	72	69
5	504 Boyz	Wobble Wobble	2000-04-15	57	34	25	17
6	98^0	Give Me Just O...	2000-08-19	51	39	34	26
7	A*Teens	Dancing Queen	2000-07-08	97	97	96	95
8	Aaliyah	I Don't Wanna	2000-01-29	84	62	51	41
9	Aaliyah	Try Again	2000-03-18	59	53	38	28
10	Adams, Yolanda	Open My Heart	2000-08-26	76	76	74	69

```
# i 307 more rows
```

Is this data tidy?

More tidy vs untidy

Is the following data tidy?

List of 3

\$:List of 8

```
..$ name      : chr "Luke Skywalker"
..$ height    : chr "172"
..$ mass      : chr "77"
..$ hair_color: chr "blond"
..$ skin_color: chr "fair"
..$ eye_color  : chr "blue"
..$ birth_year: chr "19BBY"
..$ gender     : chr "male"
```

\$:List of 8

```
..$ name      : chr "C-3P0"
..$ height    : chr "167"
..$ mass      : chr "75"
..$ hair_color: chr "n/a"
..$ skin_color: chr "gold"
..$ eye_color  : chr "yellow"
```

List of 3

\$:List of 8

```
..$ name      : chr "Darth Vader"
..$ height    : chr "202"
..$ mass      : chr "136"
..$ hair_color: chr "none"
..$ skin_color: chr "white"
..$ eye_color  : chr "yellow"
..$ birth_year: chr "41.9BBY"
..$ gender     : chr "male"
```

\$:List of 8

```
..$ name      : chr "Leia Organa"
..$ height    : chr "150"
..$ mass      : chr "49"
..$ hair_color: chr "brown"
..$ skin_color: chr "light"
..$ eye_color  : chr "brown"
```



Modern data frames

The tidyverse includes the tibble package that extends data frames to be a bit more modern. The core features of tibbles is to have a nicer printing method as well as being “surly” and “lazy”.

```
1 library(tibble)
```



```
1 iris
```

	Sepal.Length	Sepal.Width	Petal.Length
1	5.1	3.5	1.4
2	4.9	3.0	1.4
3	4.7	3.2	1.3
4	4.6	3.1	1.5
5	5.0	3.6	1.4
6	5.4	3.9	1.7
7	4.6	3.4	1.4
8	5.0	3.4	1.5
9	4.4	2.9	1.4
10	4.9	3.1	1.5
11	5.4	3.7	1.5
12	4.8	3.4	1.6
13	4.8	3.0	1.4
14	4.3	3.0	1.1
15	5.8	4.0	1.2
16	5.7	4.4	1.5
17	5.4	3.9	1.3
18	5.1	3.5	1.4
19	5.7	3.8	1.7
20	5.1	3.8	1.5
21	5.4	3.4	1.7
22	5.1	3.7	1.5

```
1 (tbl_iris = as_tibble(iris))
```

```
# A tibble: 150 × 5
  Sepal.Length Sepal.Width Petal.Length
      <dbl>         <dbl>         <dbl>
1         5.1         3.5         1.4
2         4.9         3.0         1.4
3         4.7         3.2         1.3
4         4.6         3.1         1.5
5          5.0         3.6         1.4
6         5.4         3.9         1.7
7         4.6         3.4         1.4
8          5.0         3.4         1.5
9         4.4         2.9         1.4
10        4.9         3.1         1.5
# i 140 more rows
# i 2 more variables: Petal.Width <dbl>,
#   Species <fct>
```

Tibbles are lazy (preserving type)

By default, subsetting tibbles always results in another tibble (`$` or `[[` can still be used to subset for a specific column). i.e. tibble subsets are always preserving and therefore type consistent.

```
1 tbl_iris[1,]
```

```
# A tibble: 1 × 5
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
	<dbl>	<dbl>	<dbl>	<dbl>	<fct>
1	5.1	3.5	1.4	0.2	setosa

```
1 tbl_iris[,1]
```

```
# A tibble: 150 × 1
  Sepal.Length
    <dbl>
1       5.1
2       4.9
3       4.7
4       4.6
5        5
6       5.4
7       4.6
8        5
9       4.4
10      4.9
# i 140 more rows
```

```
1 head(tbl_iris[[1]])
```

```
[1] 5.1 4.9 4.7 4.6 5.0 5.4
```

```
1 head(tbl_iris$Species)
```

```
[1] setosa setosa setosa setosa setosa set
Levels: setosa versicolor virginica
```

Tibbles are lazy (partial matching)

Tibbles do not use partial matching when the `$` operator is used.

```
1 head( iris$Species )
```

```
[1] setosa setosa setosa setosa  
Levels: setosa versicolor virg
```

```
1 head( tbl_iris$Species )
```

```
[1] setosa setosa setosa setosa  
Levels: setosa versicolor virg
```

```
1 head( iris$Sp )
```

```
[1] setosa setosa setosa setosa  
Levels: setosa versicolor virg
```

```
1 head( tbl_iris$Sp )
```

```
Warning: Unknown or uninitialised  
NULL
```

Tibbles are lazy (length coercion)

Only vectors with length 1 will undergo length coercion / recycling - anything else throws an error.

```
1 data.frame(x = 1:4, y = 1)
```

	x	y
1	1	1
2	2	1
3	3	1
4	4	1

```
1 tibble(x = 1:4, y = 1)
```

```
# A tibble: 4 × 2
      x     y
  <int> <dbl>
1     1     1
2     2     1
3     3     1
4     4     1
```

```
1 data.frame(x = 1:4, y = 1:2)
```

	x	y
1	1	1
2	2	2
3	3	1
4	4	2

```
1 tibble(x = 1:4, y = 1:2)
```

```
Error in `tibble()`:
! Tibble columns must have compatible sizes
• Size 4: Existing data.
• Size 2: Column `y`.
i Only values of size one are recycled
```

Tibbles and S3

```
1 t = tibble(  
2   x = 1:3,  
3   y = c("A","B","C")  
4 )  
5  
6 class(t)
```

```
[1] "tbl_df"      "tbl"        "data.frame" [1] "data.frame"
```

```
1 d = data.frame(  
2   x = 1:3,  
3   y = c("A","B","C")  
4 )  
5  
6 class(d)
```

```
1 methods(class="tbl_df")
```

```
[1] [      [[      [[<-      [<-      $  
[6] $<-      as.data.frame coerce      initialize names<-  
[11] Ops      row.names<-  show      slotsFromS3 str  
[16] tbl_sum  
see '?methods' for accessing help and source code
```

```
1 methods(class="tbl")
```

```
[1] [[<-      [<-      $<-      coerce      format  
[6] glimpse      initialize Ops      print      show  
[11] slotsFromS3 tbl_sum  
see '?methods' for accessing help and source code
```

Tibble support?

Tibbles are just specialized data frames, and will fall back to base data frame methods when needed.

```
1 d = tibble(  
2   x = rnorm(100),  
3   y = 3 + x + rnorm(100, sd = 0.1)  
4 )
```

```
1 lm(y~x, data = d)
```

Call:

```
lm(formula = y ~ x, data = d)
```

Coefficients:

(Intercept)	x
2.996	1.005

Why did this work?



magrittr

Sta 523 - Fall 2025

What is a pipe

In software engineering, a pipeline consists of a chain of processing elements (processes, threads, coroutines, functions, etc.), arranged so that the output of each element is the input of the next;

[Wikipedia - Pipeline \(software\)](#)

Magrittr's pipe is an infix operator that allows us to link two functions together in a way that is readable from left to right.

The two code examples below are equivalent,

```
1 f(g(x=1, y=2), n=2)
```

```
1 g(x=1, y=2) %>% f(n=2)
```

Readability

Consider the following sequence of actions that describe the process of getting to campus in the morning:

I need to find my key, then unlock my car, then start my car, then drive to school, then park.

Expressed as a set of nested functions in R pseudocode this would look like:

```
1 park(drive(start_car(find("keys")), to="campus"))
```

Writing it out using pipes give it a more natural (and easier to read) structure:

```
1 find("keys") %>%  
2   start_car() %>%  
3   drive(to="campus") %>%  
4   park()
```

Approaches

All of the following are fine, it comes down to personal preference:

Nested:

```
1 h( g( f(x), y=1), z=1 )
```

Piped:

```
1 f(x) %>%  
2   g(y=1) %>%  
3   h(z=1)
```

Intermediate:

```
1 res = f(x)  
2 res = g(res, y=1)  
3 res = h(res, z=1)
```

What about other arguments?

Sometimes we want to send our results to an function argument other than first one or we want to use the previous result for multiple arguments. In these cases we can refer to the previous result using `..`

```
1 data.frame(a = 1:3, b = 3:1) %>% lm(a~b, data=.)
```

Call:

```
lm(formula = a ~ b, data = .)
```

Coefficients:

(Intercept)	b
4	-1

```
1 data.frame(a = 1:3, b = 3:1) %>% .[[1]]
```

```
[1] 1 2 3
```

```
1 data.frame(a = 1:3, b = 3:1) %>% .[[length(.)]]
```

```
[1] 3 2 1
```

The base R pipe

As of R v4.1.0 a native pipe operator was added to the base language in R, it is implemented as `|>`.

```
1 1:10 |> cumsum()
```

```
[1] 1 3 6 10 15 21 28 36 45 55
```

```
1 1:10 |> cumsum() |> mean()
```

```
[1] 22
```


The current version of RStudio on the departmental servers is v4.4.1 so you are welcome to use it.

Base R pipe considerations:

- Depending on an R version ≥ 4.1 is a harder dependency than depending on the `magrittr` package
- `|>` has less overhead than `%>%` but the difference is unlikely to matter in practice most of the time
- `|>` supports an equivalent to `.` using `_` as of R v4.2 (but only for named arguments)

```
1 data.frame(a = 1:3, b = 3:1) |>  
2   lm(a~b, data=_)
```

Generally we will prefer the base pipe in this class, but using either is fine.



A Grammar of Data Manipulation

dplyr is based on the concepts of functions as verbs that manipulate data frames.

Core single data frame functions / verbs:

- `filter()` / `slice()` - pick rows based on criteria
- `select()` / `rename()` - select columns by name
- `pull()` - grab a column as a vector
- `arrange()` - reorder rows
- `mutate()` / `transmute()` - create or modify columns
- `distinct()` - filter for unique rows
- `summarise()` / `count()` - reduce variables to values
- `group_by()` / `ungroup()` - modify other verbs to act on subsets
- `relocate()` - change column order
- ... (many more)

dplyr rules

1. First argument is *always* a data frame
2. Subsequent arguments say what to do with the data frame
3. *Always* return a data frame
4. Don't modify in place
5. Magic via non-standard evaluation + lazy evaluation and S3

Example Data

We will demonstrate dplyr's functionality using the nycflights13 data.

```
1 library(dplyr)
2 library(nycflights13)
```

```
1 flights
```

```
# A tibble: 336,776 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

```
# i 336,766 more rows
```

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>
```

filter() - March flights

```
1 flights |> filter(month == 3)
```

```
# A tibble: 28,834 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	3	1	4	2159	125	318
2	2013	3	1	50	2358	52	526
3	2013	3	1	117	2245	152	223
4	2013	3	1	454	500	-6	633
5	2013	3	1	505	515	-10	746
6	2013	3	1	521	530	-9	813
7	2013	3	1	537	540	-3	856
8	2013	3	1	541	545	-4	1014
9	2013	3	1	549	600	-11	639
10	2013	3	1	550	600	-10	747

```
# i 28,824 more rows
```

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
```

filter() - Flights in the first 7 days of March

```
1 flights |> filter(month == 3, day <= 7)
```

```
# A tibble: 6,530 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	3	1	4	2159	125	318
2	2013	3	1	50	2358	52	526
3	2013	3	1	117	2245	152	223
4	2013	3	1	454	500	-6	633
5	2013	3	1	505	515	-10	746
6	2013	3	1	521	530	-9	813
7	2013	3	1	537	540	-3	856
8	2013	3	1	541	545	-4	1014
9	2013	3	1	549	600	-11	639
10	2013	3	1	550	600	-10	747

```
# i 6,520 more rows
```

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
```

filter() - Flights to LAX or JFK in March

```
1 flights |> filter(dest == "LAX" | dest == "JFK", month==3)
```

```
# A tibble: 1,178 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	3	1	607	610	-3	832
2	2013	3	1	629	632	-3	844
3	2013	3	1	657	700	-3	953
4	2013	3	1	714	715	-1	939
5	2013	3	1	716	710	6	958
6	2013	3	1	727	730	-3	1007
7	2013	3	1	836	840	-4	1111
8	2013	3	1	857	900	-3	1202
9	2013	3	1	903	900	3	1157
10	2013	3	1	904	831	33	1150

```
# i 1,168 more rows
```

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
```

slice() - First 10 flights

```
1 flights |> slice(1:10)
```

```
# A tibble: 10 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
# carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest
```


slice() - Last 5 flights

```
1 flights |> slice((n()-4):n())
```

```
# A tibble: 5 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	9	30	NA	1455	NA	NA
2	2013	9	30	NA	2200	NA	NA
3	2013	9	30	NA	1210	NA	NA
4	2013	9	30	NA	1159	NA	NA
5	2013	9	30	NA	840	NA	NA

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
# carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest  
# air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
# time_hour <dtm>
```

slice_tail() - Last 5 flights

```
1 flights |> slice_tail(n = 5)
```

```
# A tibble: 5 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	9	30	NA	1455	NA	NA
2	2013	9	30	NA	2200	NA	NA
3	2013	9	30	NA	1210	NA	NA
4	2013	9	30	NA	1159	NA	NA
5	2013	9	30	NA	840	NA	NA

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
# carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest  
# air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
# time_hour <dtm>
```

select() - Individual Columns

```
1 flights |> select(year, month, day)
```

```
# A tibble: 336,776 × 3
```

	year	month	day
	<int>	<int>	<int>
1	2013	1	1
2	2013	1	1
3	2013	1	1
4	2013	1	1
5	2013	1	1
6	2013	1	1
7	2013	1	1
8	2013	1	1
9	2013	1	1
10	2013	1	1

```
# i 336,766 more rows
```

select() - Exclude Columns

```
1 flights |> select(-year, -month, -day)
```

```
# A tibble: 336,776 × 16
```

	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay
	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	517	515	2	830	819	11
2	533	529	4	850	830	20
3	542	540	2	923	850	33
4	544	545	-1	1004	1022	-18
5	554	600	-6	812	837	-25
6	554	558	-4	740	728	12
7	555	600	-5	913	854	19
8	557	600	-3	709	723	-14
9	557	600	-3	838	846	-8
10	558	600	-2	753	745	8

```
# i 336,766 more rows
```

```
# i 10 more variables: carrier <chr>, flight <int>, tailnum <chr>,  
# origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
# minute <dbl>, time_hour <dttm>
```

select() - Ranges

```
1 flights |> select(year:day)
```

```
# A tibble: 336,776 × 3
```

	year	month	day
	<int>	<int>	<int>
1	2013	1	1
2	2013	1	1
3	2013	1	1
4	2013	1	1
5	2013	1	1
6	2013	1	1
7	2013	1	1
8	2013	1	1
9	2013	1	1
10	2013	1	1

```
# i 336,766 more rows
```

select() - Exclusion Ranges

```
1 flights |> select(-(year:day))
```

```
# A tibble: 336,776 × 16
```

	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay
	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	517	515	2	830	819	11
2	533	529	4	850	830	20
3	542	540	2	923	850	33
4	544	545	-1	1004	1022	-18
5	554	600	-6	812	837	-25
6	554	558	-4	740	728	12
7	555	600	-5	913	854	19
8	557	600	-3	709	723	-14
9	557	600	-3	838	846	-8
10	558	600	-2	753	745	8

```
# i 336,766 more rows
```

```
# i 10 more variables: carrier <chr>, flight <int>, tailnum <chr>,  
# origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
# minute <dbl>, time_hour <dtm>
```

select() - Matching contains()

```
1 flights |> select(contains("dep"), contains("arr"))
```

```
# A tibble: 336,776 × 7
```

	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay
	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	517	515	2	830	819	11
2	533	529	4	850	830	20
3	542	540	2	923	850	33
4	544	545	-1	1004	1022	-18
5	554	600	-6	812	837	-25
6	554	558	-4	740	728	12
7	555	600	-5	913	854	19
8	557	600	-3	709	723	-14
9	557	600	-3	838	846	-8
10	558	600	-2	753	745	8

```
# i 336,766 more rows
```

```
# i 1 more variable: carrier <chr>
```

select() - Matching starts_with()

```
1 flights |> select(starts_with("dep"), starts_with("arr"))
```

```
# A tibble: 336,776 × 4
```

	dep_time	dep_delay	arr_time	arr_delay
	<int>	<dbl>	<int>	<dbl>
1	517	2	830	11
2	533	4	850	20
3	542	2	923	33
4	544	-1	1004	-18
5	554	-6	812	-25
6	554	-4	740	12
7	555	-5	913	19
8	557	-3	709	-14
9	557	-3	838	-8
10	558	-2	753	8

```
# i 336,766 more rows
```

Other helpers provide by [tidyselect](#):

select() + where() - Get numeric columns

```
1 flights |> select(where(is.numeric))
```

```
# A tibble: 336,776 × 14
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

```
# i 336,766 more rows
```

```
# i 7 more variables: sched_arr_time <int>, arr_delay <dbl>, flight <int>
```

```
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>
```

select() + where() - Get non-numeric columns

```
1 flights |> select(where(function(x) !is.numeric(x)))
```

```
# A tibble: 336,776 × 5
```

	carrier	tailnum	origin	dest	time_hour
	<chr>	<chr>	<chr>	<chr>	<dtm>
1	UA	N14228	EWR	IAH	2013-01-01 05:00:00
2	UA	N24211	LGA	IAH	2013-01-01 05:00:00
3	AA	N619AA	JFK	MIA	2013-01-01 05:00:00
4	B6	N804JB	JFK	BQN	2013-01-01 05:00:00
5	DL	N668DN	LGA	ATL	2013-01-01 06:00:00
6	UA	N39463	EWR	ORD	2013-01-01 05:00:00
7	B6	N516JB	EWR	FLL	2013-01-01 06:00:00
8	EV	N829AS	LGA	IAD	2013-01-01 06:00:00
9	B6	N593JB	JFK	MCO	2013-01-01 06:00:00
10	AA	N3ALAA	LGA	ORD	2013-01-01 06:00:00

```
# i 336,766 more rows
```

relocate - to the front

```
1 flights |> relocate(carrier, origin, dest)
```

```
# A tibble: 336,776 × 19
```

	carrier	origin	dest	year	month	day	dep_time	sched_dep_time	dep_delay
	<chr>	<chr>	<chr>	<int>	<int>	<int>	<int>	<int>	<dbl>
1	UA	EWR	IAH	2013	1	1	517	515	2
2	UA	LGA	IAH	2013	1	1	533	529	4
3	AA	JFK	MIA	2013	1	1	542	540	2
4	B6	JFK	BQN	2013	1	1	544	545	-1
5	DL	LGA	ATL	2013	1	1	554	600	-6
6	UA	EWR	ORD	2013	1	1	554	558	-4
7	B6	EWR	FLL	2013	1	1	555	600	-5
8	EV	LGA	IAD	2013	1	1	557	600	-3
9	B6	JFK	MCO	2013	1	1	557	600	-3
10	AA	LGA	ORD	2013	1	1	558	600	-2

```
# i 336,766 more rows
```

```
# i 10 more variables: arr_time <int>, sched_arr_time <int>,  
#   arr_delay <dbl>, flight <int>, tailnum <chr>, air_time <dbl>,  
#   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

relocate - to the end

```
1 flights |> relocate(year, month, day, .after = last_col())
```

```
# A tibble: 336,776 × 19
```

	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay
	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	517	515	2	830	819	11
2	533	529	4	850	830	20
3	542	540	2	923	850	33
4	544	545	-1	1004	1022	-18
5	554	600	-6	812	837	-25
6	554	558	-4	740	728	12
7	555	600	-5	913	854	19
8	557	600	-3	709	723	-14
9	557	600	-3	838	846	-8
10	558	600	-2	753	745	8

```
# i 336,766 more rows
```

```
# i 13 more variables: carrier <chr>, flight <int>, tailnum <chr>,  
# origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
# minute <dbl>, time_hour <dtm>, year <int>, month <int>, day <int>
```

rename() - Change column names

```
1 flights |> rename(tail_number = tailnum)
```

```
# A tibble: 336,776 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

```
# i 336,766 more rows  
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tail_number <chr>, origin <chr>,  
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#   time_hour <dtm>
```

select() vs. rename()

```
1 flights |> select(tail_number = tailnum)
```

```
# A tibble: 336,776 × 1
```

```
  tail_number  
  <chr>  
1 N14228  
2 N24211  
3 N619AA  
4 N804JB  
5 N668DN  
6 N39463  
7 N516JB  
8 N829AS  
9 N593JB  
10 N3ALAA  
# i 336,766 more rows
```

```
1 flights |> rename(tail_number = tailnum)
```

```
# A tibble: 336,776 × 19
```

```
  year month   day dep_time sched_dep_time  
  <int> <int> <int>   <int>         <int>  
1  2013     1     1     517           515  
2  2013     1     1     533           529  
3  2013     1     1     542           540  
4  2013     1     1     544           545  
5  2013     1     1     554           600  
6  2013     1     1     554           558  
7  2013     1     1     555           600  
8  2013     1     1     557           600  
9  2013     1     1     557           600  
10 2013     1     1     558           600  
# i 336,766 more rows  
# i 14 more variables: dep_delay <dbl>,  
#   arr_time <int>, sched_arr_time <int>,  
#   arr_delay <dbl>, carrier <chr>,  
#   flight <int>, tail_number <chr>,  
#   origin <chr>, dest <chr>,  
#   air_time <dbl>, distance <dbl>, ...
```

pull()

```
1 names(flights)
```

```
[1] "year"          "month"          "day"
[4] "dep_time"      "sched_dep_time" "dep_delay"
[7] "arr_time"      "sched_arr_time" "arr_delay"
[10] "carrier"       "flight"         "tailnum"
[13] "origin"        "dest"           "air_time"
[16] "distance"      "hour"           "minute"
[19] "time_hour"
```

```
1 flights |> pull("year") |> head()
```

```
[1] 2013 2013 2013 2013 2013 2013
```

```
1 flights |> pull(1) |> head()
```

```
[1] 2013 2013 2013 2013 2013 2013
```

```
1 flights |> pull(-1) |> head()
```

```
[1] "2013-01-01 05:00:00 EST" "2013-01-01 05:00:00 EST"
[3] "2013-01-01 05:00:00 EST" "2013-01-01 05:00:00 EST"
[5] "2013-01-01 06:00:00 EST" "2013-01-01 05:00:00 EST"
```

arrange() - Sort data

```
1 flights |>
2   filter(month==3,day==2) |>
3   arrange(origin, dest)
```

A tibble: 765 × 19

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	3	2	1336	1329	7	1426
2	2013	3	2	628	629	-1	837
3	2013	3	2	637	640	-3	903
4	2013	3	2	743	745	-2	945
5	2013	3	2	857	900	-3	1117
6	2013	3	2	1027	1030	-3	1234
7	2013	3	2	1134	1145	-11	1332
8	2013	3	2	1412	1415	-3	1636
9	2013	3	2	1633	1636	-3	1848
10	2013	3	2	1655	1700	-5	1857

i 755 more rows

i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
minute <dbl>, time_hour <dtm>

arrange() w/ desc() - descending order

```
1 flights |>
2   filter(month==3, day==2) |>
3   arrange(desc(origin), dest) |>
4   select(origin, dest, tailnum)
```

```
# A tibble: 765 × 3
  origin dest  tailnum
  <chr>   <chr> <chr>
1 LGA     ATL   N928AT
2 LGA     ATL   N623DL
3 LGA     ATL   N680DA
4 LGA     ATL   N996AT
5 LGA     ATL   N510MQ
6 LGA     ATL   N663DN
7 LGA     ATL   N942DL
8 LGA     ATL   N511MQ
9 LGA     ATL   N910DE
10 LGA    ATL   N902DE
# i 755 more rows
```

distinct() - Find unique rows

```
1 flights |>
2   select(origin, dest) |>
3   distinct() |>
4   arrange(origin,dest)
```

```
# A tibble: 224 × 2
```

```
  origin dest
  <chr>  <chr>
1 EWR    ALB
2 EWR    ANC
3 EWR    ATL
4 EWR    AUS
5 EWR    AVL
6 EWR    BDL
7 EWR    BNA
8 EWR    BOS
9 EWR    BQN
10 EWR    BTV
```

```
# i 214 more rows
```

mutate() - Modify / create columns

```
1 flights |>
2   select(year:day) |>
3   mutate(date = paste(year, month, day, sep="/"))
```

```
# A tibble: 336,776 × 4
  year month   day date
  <int> <int> <int> <chr>
1  2013     1     1 2013/1/1
2  2013     1     1 2013/1/1
3  2013     1     1 2013/1/1
4  2013     1     1 2013/1/1
5  2013     1     1 2013/1/1
6  2013     1     1 2013/1/1
7  2013     1     1 2013/1/1
8  2013     1     1 2013/1/1
9  2013     1     1 2013/1/1
10 2013     1     1 2013/1/1
# i 336,766 more rows
```

summarise() - Aggregate rows

```
1 flights |>
2   summarize(n(), min(dep_delay), max(dep_delay))
```

```
# A tibble: 1 × 3
  `n()` `min(dep_delay)` `max(dep_delay)`
  <int>         <dbl>         <dbl>
1 336776          NA          NA
```

```
1 flights |>
2   summarize(
3     n = n(),
4     min_dep_delay = min(dep_delay, na.rm = TRUE),
5     max_dep_delay = max(dep_delay, na.rm = TRUE)
6   )
```

```
# A tibble: 1 × 3
  n min_dep_delay max_dep_delay
  <int>         <dbl>         <dbl>
1 336776        -43         1301
```

group_by()

```
1 flights |> group_by(origin)
```

```
# A tibble: 336,776 × 19
```

```
# Groups:   origin [3]
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

```
# i 336,766 more rows
```

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
```

```
# carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
```

```
# dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
# minute <dbl>, time_hour <dtm>
```

summarise() with group_by()

```
1 flights |>
2   group_by(origin) |>
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm = TRUE),
6     max_dep_delay = max(dep_delay, na.rm = TRUE)
7   )
```

A tibble: 3 × 4

	origin	n	min_dep_delay	max_dep_delay
	<chr>	<int>	<dbl>	<dbl>
1	EWR	120835	-25	1126
2	JFK	111279	-43	1301
3	LGA	104662	-33	911

Groups after summarise

```
1 flights |>
2   group_by(origin, month) |>
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm=TRUE),
6     max_dep_delay = max(dep_delay, na.rm=TRUE)
7   )
```

`summarise()` has grouped output by 'origin'. You can override using the `.groups` argument.

A tibble: 36 × 5

Groups: origin [3]

	origin	month	n	min_dep_delay	max_dep_delay
	<chr>	<int>	<int>	<dbl>	<dbl>
1	EWR	1	9893	-21	1126
2	EWR	2	9107	-21	786
3	EWR	3	10420	-22	443
4	EWR	4	10531	-21	545
5	EWR	5	10592	-20	878
6	EWR	6	10175	-19	502
7	EWR	7	10475	-18	653
8	EWR	8	10359	-17	424
9	EWR	9	9550	-23	486

Avoid the message

```
1 flights |>
2   group_by(origin, month) |>
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm=TRUE),
6     max_dep_delay = max(dep_delay, na.rm=TRUE),
7     .groups = "drop"
8   )
```

A tibble: 36 × 5

	origin	month	n	min_dep_delay
	<chr>	<int>	<int>	<dbl>
1	EWR	1	9893	-21
2	EWR	2	9107	-21
3	EWR	3	10420	-22
4	EWR	4	10531	-21
5	EWR	5	10592	-20
6	EWR	6	10175	-19
7	EWR	7	10475	-18
8	EWR	8	10359	-17
9	EWR	9	9550	-23
10	EWR	10	10104	-25

i 26 more rows

i 1 more variable: max_dep_delay <dbl>

```
1 flights |>
2   group_by(origin, month) |>
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm=TRUE),
6     max_dep_delay = max(dep_delay, na.rm=TRUE),
7     .groups = "keep"
8   )
```

A tibble: 36 × 5

Groups: origin, month [36]

	origin	month	n	min_dep_delay
	<chr>	<int>	<int>	<dbl>
1	EWR	1	9893	-21
2	EWR	2	9107	-21
3	EWR	3	10420	-22
4	EWR	4	10531	-21
5	EWR	5	10592	-20
6	EWR	6	10175	-19
7	EWR	7	10475	-18
8	EWR	8	10359	-17
9	EWR	9	9550	-23
10	EWR	10	10104	-25

i 26 more rows

i 1 more variable: max_dep_delay <dbl>

The .by argument

The `.by` (and `by`) arguments are used for per operation grouping while `group_by()` is intended for persistent grouping. See `?dplyr_by` for more details and examples.

```
1 flights |>
2   summarize(
3     n = n(),
4     min_dep_delay = min(dep_delay, na.rm=TRUE),
5     max_dep_delay = max(dep_delay, na.rm=TRUE),
6     .by = origin
7   )
```

```
# A tibble: 3 × 4
  origin      n min_dep_delay max_dep_delay
  <chr>    <int>         <dbl>         <dbl>
1 EWR    120835          -25          1126
2 LGA    104662          -33           911
3 JFK    111279          -43          1301
```

count()

```
1 flights |>
2   summarize(
3     n = n(),
4     .by = c(origin, carrier)
5   )
```

```
# A tibble: 35 × 3
  origin carrier      n
  <chr>   <chr> <int>
1 EWR     UA    46087
2 LGA     UA     8044
3 JFK     AA    13783
4 JFK     B6    42076
5 LGA     DL    23067
6 EWR     B6     6557
7 LGA     EV     8826
8 LGA     AA    15459
9 JFK     UA     4534
10 LGA     B6     6002
# i 25 more rows
```

```
1 flights |>
2   count(origin, carrier)
```

```
# A tibble: 35 × 3
  origin carrier      n
  <chr>   <chr> <int>
1 EWR     9E    1268
2 EWR     AA   3487
3 EWR     AS     714
4 EWR     B6   6557
5 EWR     DL   4342
6 EWR     EV  43939
7 EWR     MQ   2276
8 EWR     00        6
9 EWR     UA   46087
10 EWR     US   4405
# i 25 more rows
```

mutate() with .by

```
1 flights |>
2   mutate(n = n(), .by = origin) |>
3   select(origin, n)
```

A tibble: 336,776 × 2

	origin	n
	<chr>	<int>
1	EWR	120835
2	LGA	104662
3	JFK	111279
4	JFK	111279
5	LGA	104662
6	EWR	120835
7	EWR	120835
8	LGA	104662
9	JFK	111279
10	LGA	104662

i 336,766 more rows

Exercises / Examples

1. How many flights to Los Angeles (LAX) did each of the legacy carriers (AA, UA, DL or US) have in May from JFK, and what was their average duration?
2. What was the shortest flight out of each airport in terms of distance? In terms of duration?
3. Which plane (check the tail number) flew out of each New York airport the most?
4. Which date should you fly on if you want to have the lowest possible average departure delay? What about arrival delay?